Proof of Proper Table Implementation:

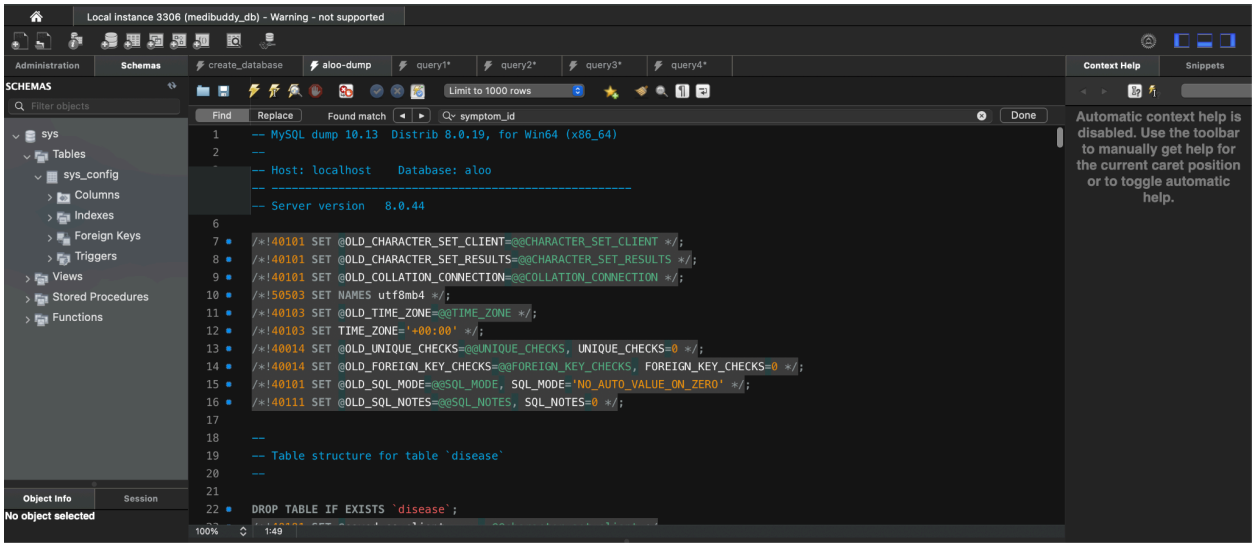

Procedure:

1. Create the database on MySQL.
2. Create the relations.
3. Find the estimated total cost of each of the 4 queries without any indices added.
4. Through forward selection, add an index for each attribute, one at a time.
5. Keep attribute indices that reduce costs.

Query 1

List of non-primary key attributes used in JOIN, WHERE, GROUP BY, and HAVING clauses:

- doctor(city)
- doctor(specialty_id)
- disease(specialty_id)
- disease_symptom(disease_id)
- disease_symptom(symptom_id)

Idx Added	Estimated Total Cost
None	19,815
doctor(city)	21,166
doctor(specialty_id)	19,123
doctor(specialty_id) + disease(specialty_id)	19,123
doctor(specialty_id) + disease(specialty_id) + disease_symptom(disease_id)	19,123

doctor(specialty_id) + disease(specialty_id) + disease_symptom(symptom_id)	13,311

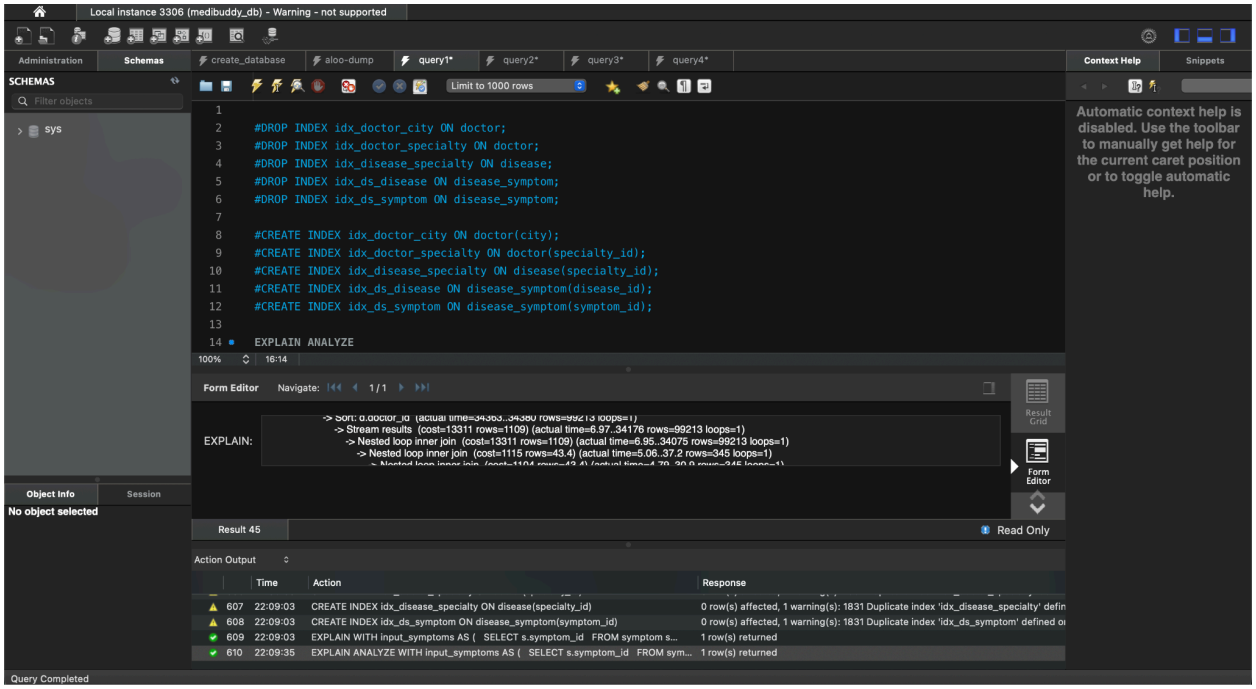

Conclusion: Through forward selection, doctor(specialty_id), disease(specialty_id), and disease(symptom_id) should have indices associated with them to minimize estimated total cost.

Query 2

List of non-primary key attributes used in JOIN, WHERE, GROUP BY, and HAVING clauses:

- doctor(city)
- doctor(specialty_id)
- disease(specialty_id)
- disease_symptom(disease_id)
- disease_symptom(symptom_id)

Idx Added	Estimated Total Cost
None	660,382
doctor(city)	187,865
doctor(city) + doctor(specialty_id)	660,382
doctor(city) + disease(specialty_id)	202,553
doctor(city) + disease_symptom(disease_id)	202,553

doctor(city) + disease_symptom(symptom_id)	202,553

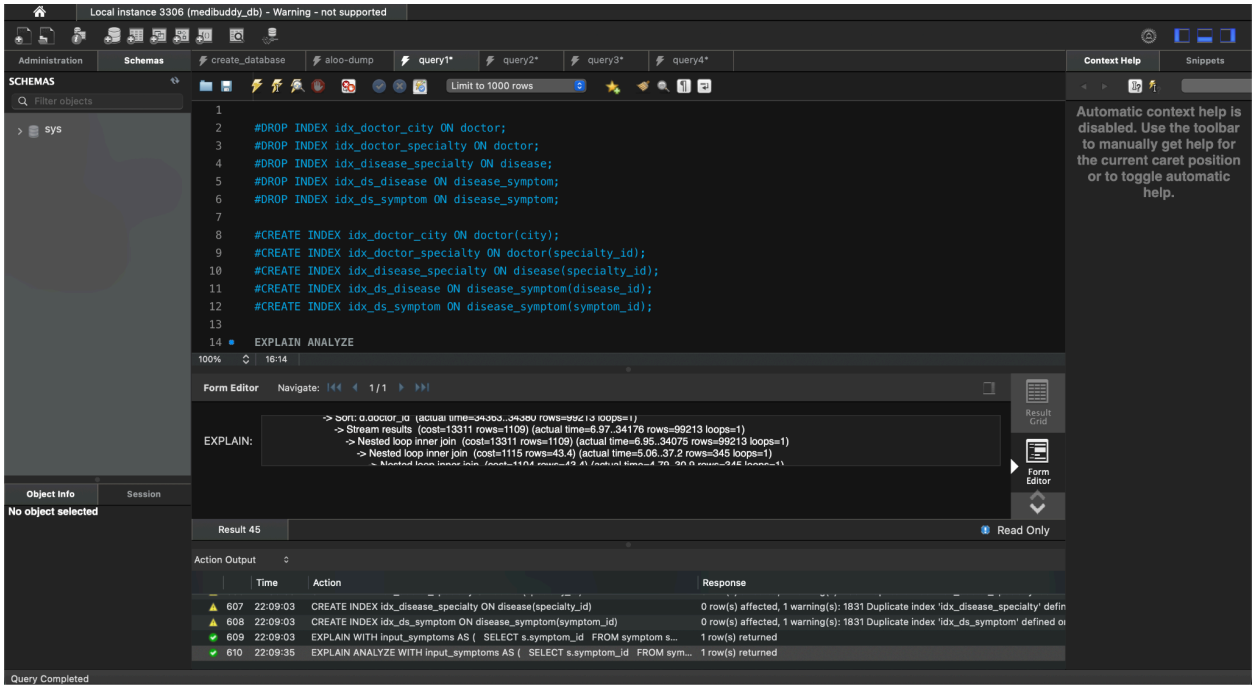

Conclusion: Through forward selection, doctor(city) should be the only attribute w/ an index associated w/ it to minimize estimated total costs.

Query 3

List of non-primary key attributes used in JOIN, WHERE, GROUP BY, and HAVING clauses:

- doctor(state)
- doctor(specialty_id)
- disease(specialty_id)
- disease_symptom(disease_id)

Idx Added	Estimated Total Cost
None	9.07e+6
doctor(state)	400e+6
doctor(specialty_id)	9.07e+6
disease(specialty_id)	9.07e+6
disease_symptom(disease_id)	8.84e+6

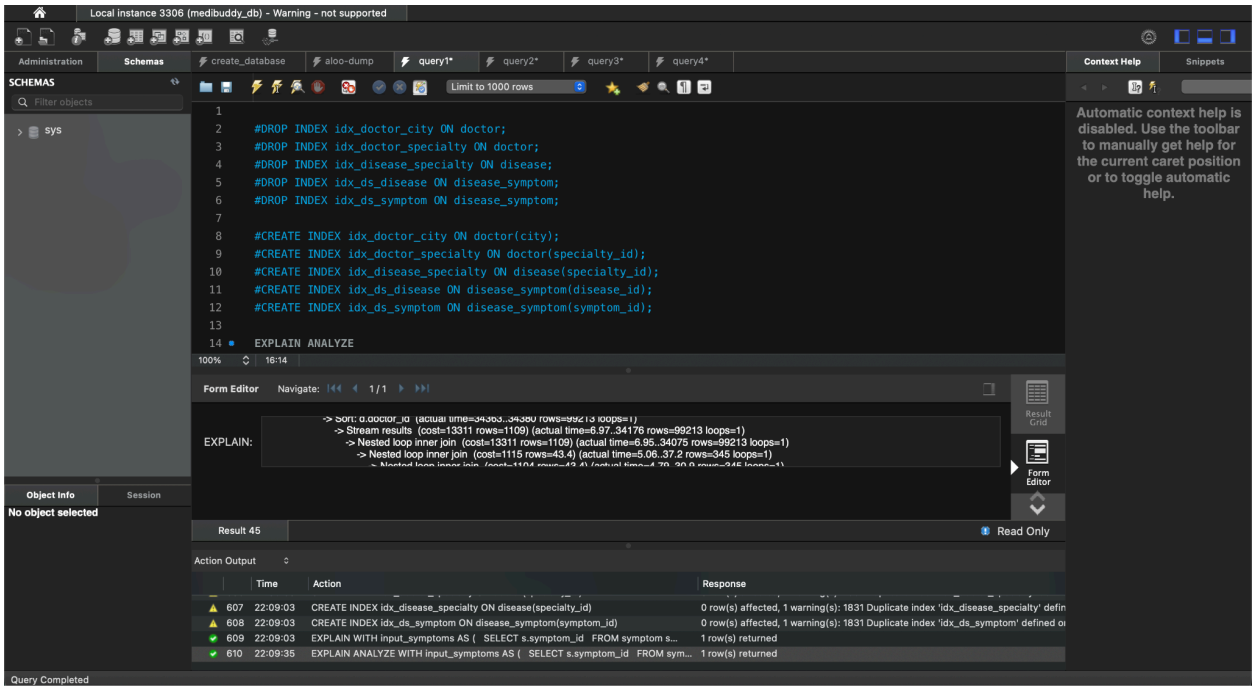

Conclusion: Through forward selection, no attributes should have an index associated w/ it to minimize estimated total costs. This could be because of factors such as doctor(state) not having many unique values, which increases lookup time.

Query 4

List of non-primary key attributes used in JOIN, WHERE, GROUP BY, and HAVING clauses:

- doctor(city)
- doctor(specialty_id)
- disease(specialty_id)

Idx Added	Estimated Total Cost
None	429,533
doctor(city)	10,345
doctor(city) + doctor(specialty_id)	10,345
doctor(city) + disease(specialty_id)	9.07e+6

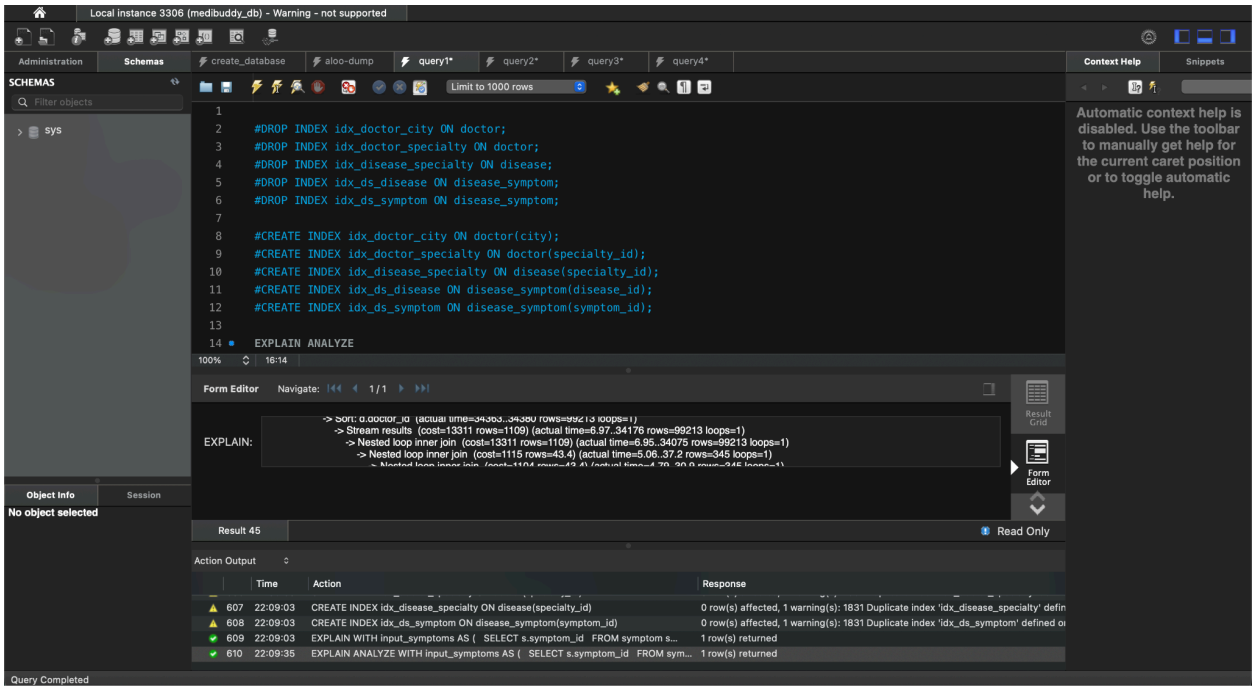

Conclusion: Through forward selection, doctor(city) should be the only attribute w/ an index associated w/ it to minimize estimated total costs.